

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv('laptop_data.csv')
```

```
In [3]: df.head()
```

Out[3]:

	Unnamed: 0	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	C
0	0	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8GB	128GB SSD	Intel Iris Plus Graphics 640	m
1	1	Apple	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8GB	128GB Flash Storage	Intel HD Graphics 6000	m
2	2	HP	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8GB	256GB SSD	Intel HD Graphics 620	N
3	3	Apple	Ultrabook	15.4	IPS Panel Retina Display 2880x1800	Intel Core i7 2.7GHz	16GB	512GB SSD	AMD Radeon Pro 455	m
4	4	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 3.1GHz	8GB	256GB SSD	Intel Iris Plus Graphics 650	m

```
In [4]: df.shape
```

Out[4]: (1303, 12)

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Unnamed: 0          1303 non-null  int64
1   Company             1303 non-null  object
2   TypeName            1303 non-null  object
3   Inches              1303 non-null  float64
4   ScreenResolution    1303 non-null  object
5   Cpu                 1303 non-null  object
6   Ram                 1303 non-null  object
7   Memory              1303 non-null  object
8   Gpu                 1303 non-null  object
9   OpSys               1303 non-null  object
10  Weight              1303 non-null  object
11  Price               1303 non-null  float64
```

dtypes: float64(2), int64(1), object(9)
memory usage: 122.3+ KB

In [6]: `df.duplicated().sum()`

Out[6]: 0

In [7]: `df.isnull().sum()`

Out[7]: Unnamed: 0 0
Company 0
TypeName 0
Inches 0
ScreenResolution 0
Cpu 0
Ram 0
Memory 0
Gpu 0
OpSys 0
Weight 0
Price 0
dtype: int64

In [8]: `df.drop(columns=['Unnamed: 0'],inplace=True)`

In [9]: `df.head()`

Out[9]:

	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	Weig
0	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8GB	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37l
1	Apple	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8GB	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34l
2	HP	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8GB	256GB SSD	Intel HD Graphics 620	No OS	1.86l
3	Apple	Ultrabook	15.4	IPS Panel Retina Display 2880x1800	Intel Core i7 2.7GHz	16GB	512GB SSD	AMD Radeon Pro 455	macOS	1.83l
4	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 3.1GHz	8GB	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37l



In [10]: `df['Ram'] = df['Ram'].str.replace('GB','')
df['Weight'] = df['Weight'].str.replace('kg','')`

In [11]: `df.head()`

Out[11]:

	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	Weight
0	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.3
1	Apple	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.3
2	HP	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.8
3	Apple	Ultrabook	15.4	IPS Panel Retina Display 2880x1800	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.8
4	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.3



In [12]:

```
df['Ram'] = df['Ram'].astype('int32')
df['Weight'] = df['Weight'].astype('float32')
```

In [13]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Company                1303 non-null   object
1   TypeName               1303 non-null   object
2   Inches                 1303 non-null   float64
3   ScreenResolution       1303 non-null   object
4   Cpu                    1303 non-null   object
5   Ram                    1303 non-null   int32
6   Memory                 1303 non-null   object
7   Gpu                    1303 non-null   object
8   OpSys                  1303 non-null   object
9   Weight                 1303 non-null   float32
10  Price                  1303 non-null   float64
dtypes: float32(1), float64(2), int32(1), object(7)
memory usage: 101.9+ KB
```

In [14]:


```
import seaborn as sns
```

In [15]:

```
sns.distplot(df['Price'])
```


```
C:\Users\Varun kumar\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)
```

```
Out[15]: <AxesSubplot:xlabel='Price', ylabel='Density'>
```




```
In [16]: df['Company'].value_counts().plot(kind='bar')
```

Out[16]: <AxesSubplot:>




```
In [17]: sns.barplot(x=df['Company'],y=df['Price'])
plt.xticks(rotation='vertical')
plt.show()
```




```
In [18]: df['TypeName'].value_counts().plot(kind='bar')
```

```
Out[18]: <AxesSubplot:>
```




```
In [19]: sns.barplot(x=df['TypeName'],y=df['Price'])
plt.xticks(rotation='vertical')
plt.show()
```



In [20]: `sns.distplot(df['Inches'])`


C:\Users\Varun kumar\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

Out[20]: <AxesSubplot:xlabel='Inches', ylabel='Density'>



In [21]: `sns.scatterplot(x=df['Inches'],y=df['Price'])`

Out[21]: <AxesSubplot:xlabel='Inches', ylabel='Price'>



In [22]: `df['ScreenResolution'].value_counts()`

```
Out[22]: Full HD 1920x1080          507
1366x768          281
IPS Panel Full HD 1920x1080        230
IPS Panel Full HD / Touchscreen 1920x1080    53
Full HD / Touchscreen 1920x1080         47
1600x900          23
Touchscreen 1366x768              16
Quad HD+ / Touchscreen 3200x1800         15
IPS Panel 4K Ultra HD 3840x2160         12
IPS Panel 4K Ultra HD / Touchscreen 3840x2160  11
4K Ultra HD / Touchscreen 3840x2160        10
IPS Panel 1366x768                7
4K Ultra HD 3840x2160              7
Touchscreen 2560x1440              7
IPS Panel Retina Display 2304x1440        6
Touchscreen 2256x1504                6
IPS Panel Quad HD+ / Touchscreen 3200x1800  6
IPS Panel Retina Display 2560x1600        6
IPS Panel Touchscreen 2560x1440          5
IPS Panel Touchscreen 1920x1200          4
IPS Panel 2560x1440                  4
IPS Panel Retina Display 2880x1800        4
1440x900                             4
IPS Panel Quad HD+ 2560x1440             3
1920x1080                             3
IPS Panel Touchscreen 1366x768           3
Touchscreen 2400x1600                   3
2560x1440                             3
Quad HD+ 3200x1800                     3
IPS Panel Full HD 2160x1440              2
IPS Panel Touchscreen / 4K Ultra HD 3840x2160  2
IPS Panel Quad HD+ 3200x1800             2
Touchscreen / Quad HD+ 3200x1800         1
IPS Panel Retina Display 2736x1824        1
IPS Panel Full HD 1920x1200              1
Touchscreen / 4K Ultra HD 3840x2160       1
Touchscreen / Full HD 1920x1080           1
IPS Panel Touchscreen 2400x1600           1
IPS Panel Full HD 1366x768               1
IPS Panel Full HD 2560x1440              1
Name: ScreenResolution, dtype: int64
```

In [23]: `df['Touchscreen'] = df['ScreenResolution'].apply(lambda x:1 if 'Touchscreen' in x else 0)`

In [24]: `df.sample(5)`

Out[24]:


	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys
549	Lenovo	Notebook	15.6	IPS Panel Full HD 1920x1080	Intel Core i7 7500U 2.7GHz	8	512GB SSD	Intel HD Graphics 620	Windows 10
1137	Dell	Notebook	15.6	1366x768	Intel Core i5 7200U 2.5GHz	8	1TB HDD	AMD Radeon R7 M445	Windows 10
1040	Lenovo	2 in 1 Convertible	15.6	IPS Panel Full HD / Touchscreen 1920x1080	Intel Core i7 6500U 2.5GHz	4	256GB SSD	Intel HD Graphics 520	Windows 10
1146	HP	Ultrabook	12.5	Full HD 1920x1080	Intel Core i7 6500U 2.50GHz	8	256GB SSD	Intel HD Graphics 520	Windows 7
455	Dell	2 in 1 Convertible	11.6	Touchscreen 1366x768	Intel Pentium Quad Core N3710 1.6GHz	4	500GB HDD	Intel HD Graphics 405	Windows 10



In [25]:

```
df['Touchscreen'].value_counts().plot(kind='bar')
```


Out[25]: <AxesSubplot:>



In [26]:

```
sns.barplot(x=df['Touchscreen'], y=df['Price'])
```

Out[26]: <AxesSubplot:xlabel='Touchscreen', ylabel='Price'>



```
In [27]: df['Ips'] = df['ScreenResolution'].apply(lambda x:1 if 'IPS' in x else 0)
```


```
In [28]: df.head()
```

Out[28]:

	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	Weigl
0	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.3
1	Apple	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.3
2	HP	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.8
3	Apple	Ultrabook	15.4	IPS Panel Retina Display 2880x1800	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.8
4	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.3


```
In [29]: df['Ips'].value_counts().plot(kind='bar')
```

Out[29]: <AxesSubplot:>



```
In [30]: sns.barplot(x=df['Ips'],y=df['Price'])
```

Out[30]: <AxesSubplot:xlabel='Ips', ylabel='Price'>



```
In [31]: new = df['ScreenResolution'].str.split('x',n=1,expand=True)
```

```
In [32]: df['X_res'] = new[0]
df['Y_res'] = new[1]
```

```
In [33]: df.sample(5)
```

	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys
432	Lenovo	Notebook	14.0	IPS Panel Full HD 1920x1080	Intel Core i3 7130U 2.7GHz	4	128GB SSD	Intel HD Graphics 620	Windows 10
1230	Lenovo	Netbook	12.5	IPS Panel Full HD / Touchscreen 1920x1080	Intel Core i7 6500U 2.50GHz	16	512GB SSD	Intel HD Graphics 520	Windows 10
787	Lenovo	Gaming	15.6	IPS Panel Full HD 1920x1080	Intel Core i7 7700HQ 2.8GHz	4	1TB HDD	Nvidia GeForce GTX 1050	Windows 10

	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys
50	Lenovo	2 in 1 Convertible	10.1	IPS Panel Touchscreen 1920x1200	Intel Atom x5-Z8550 1.44GHz	4	64GB Flash Storage	Intel HD Graphics 400	Android
1263	Acer	Notebook	15.6	1366x768	Intel Celeron Dual Core N3060 1.6GHz	4	500GB HDD	Intel HD Graphics 400	Linux



In [34]:

```
df['X_res'] = df['X_res'].str.replace(',', '').str.findall(r'(\d+\.\d+)').apply(lambda
```

In [35]:

```
df.head()
```

Out[35]:

	Company	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys	Weight
0	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.3
1	Apple	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.3
2	HP	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.8
3	Apple	Ultrabook	15.4	IPS Panel Retina Display 2880x1800	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.8
4	Apple	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.3



In [36]:

```
df['X_res'] = df['X_res'].astype('int')
df['Y_res'] = df['Y_res'].astype('int')
```

In [37]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Company              1303 non-null   object
1   TypeName              1303 non-null   object
2   Inches                1303 non-null   float64
```

```
3  ScreenResolution  1303 non-null  object
4  Cpu               1303 non-null  object
5  Ram               1303 non-null  int32
6  Memory            1303 non-null  object
7  Gpu               1303 non-null  object
8  OpSys             1303 non-null  object
9  Weight            1303 non-null  float32
10 Price             1303 non-null  float64
11 Touchscreen       1303 non-null  int64
12 Ips               1303 non-null  int64
13 X_res             1303 non-null  int32
14 Y_res             1303 non-null  int32
dtypes: float32(1), float64(2), int32(3), int64(2), object(7)
memory usage: 132.5+ KB
```

```
In [38]: df.corr()['Price']
```

```
Out[38]: Inches      0.068197
Ram          0.743007
Weight       0.210370
Price        1.000000
Touchscreen  0.191226
Ips          0.252208
X_res        0.556529
Y_res        0.552809
Name: Price, dtype: float64
```

```
In [39]: df['ppi'] = (((df['X_res']**2) + (df['Y_res']**2))*0.5/df['Inches']).astype('float')
```

```
In [40]: df.corr()['Price']
```

```
Out[40]: Inches      0.068197
Ram          0.743007
Weight       0.210370
Price        1.000000
Touchscreen  0.191226
Ips          0.252208
X_res        0.556529
Y_res        0.552809
ppi          0.473487
Name: Price, dtype: float64
```

```
In [41]: df.drop(columns=['ScreenResolution'],inplace=True)
```

```
In [42]: df.head()
```

Out[42]:

	Company	TypeName	Inches	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	T
0	Apple	Ultrabook	13.3	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	
1	Apple	Ultrabook	13.3	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34	47895.5232	
2	HP	Notebook	15.6	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.86	30636.0000	

	Company	TypeName	Inches	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	Tc
3	Apple	Ultrabook	15.4	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.83	135195.3360	
4	Apple	Ultrabook	13.3	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	

In [43]:

```
df.drop(columns=['Inches', 'X_res', 'Y_res'], inplace=True)
```

In [44]:

```
df.head()
```

Out[44]:

	Company	TypeName	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	Touchscreen
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	
1	Apple	Ultrabook	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34	47895.5232	
2	HP	Notebook	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.86	30636.0000	
3	Apple	Ultrabook	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.83	135195.3360	
4	Apple	Ultrabook	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	

In [45]:

```
df['Cpu'].value_counts()
```

Out[45]:

```
Intel Core i5 7200U 2.5GHz    190
Intel Core i7 7700HQ 2.8GHz    146
Intel Core i7 7500U 2.7GHz    134
Intel Core i7 8550U 1.8GHz     73
Intel Core i5 8250U 1.6GHz     72
...
Intel Core i5 2.0GHz          1
AMD Ryzen 1600 3.2GHz         1
AMD A6-Series 7310 2GHz       1
Intel Core i5 6440HQ 2.6GHz    1
AMD A4-Series 7210 2.2GHz     1
Name: Cpu, Length: 118, dtype: int64
```

In [46]:

```
df['Cpu Name'] = df['Cpu'].apply(lambda x: " ".join(x.split()[0:3]))
```

In [47]:

df.head()

Out[47]:

	Company	TypeName	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	Touchscreen
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	
1	Apple	Ultrabook	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34	47895.5232	
2	HP	Notebook	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.86	30636.0000	
3	Apple	Ultrabook	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.83	135195.3360	
4	Apple	Ultrabook	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	

In [48]:

```
def fetch_processor(text):
    if text == 'Intel Core i7' or text == 'Intel Core i5' or text == 'Intel Core i3':
        return text
    else:
        if text.split()[0] == 'Intel':
            return 'Other Intel Processor'
        else:
            return 'AMD Processor'
```

In [49]:

df['Cpu brand'] = df['Cpu Name'].apply(fetch_processor)

In [50]:

df.head()

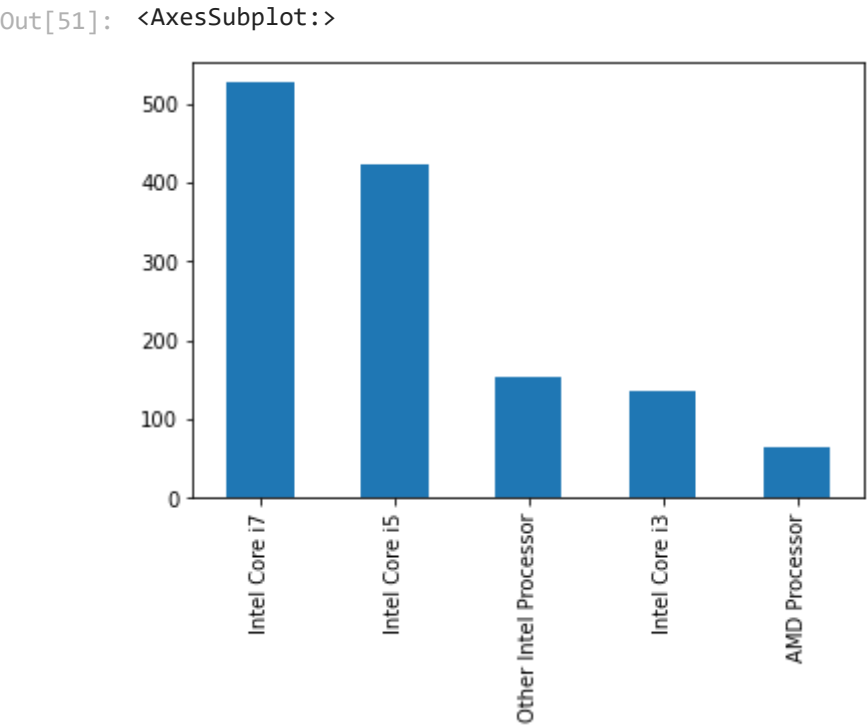
Out[50]:

	Company	TypeName	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	Touchscreen
0	Apple	Ultrabook	Intel Core i5 2.3GHz	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	
1	Apple	Ultrabook	Intel Core i5 1.8GHz	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34	47895.5232	
2	HP	Notebook	Intel Core i5 7200U 2.5GHz	8	256GB SSD	Intel HD Graphics 620	No OS	1.86	30636.0000	


	Company	TypeName	Cpu	Ram	Memory	Gpu	OpSys	Weight	Price	Touchscreen
3	Apple	Ultrabook	Intel Core i7 2.7GHz	16	512GB SSD	AMD Radeon Pro 455	macOS	1.83	135195.3360	
4	Apple	Ultrabook	Intel Core i5 3.1GHz	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	



```
In [51]: df['Cpu brand'].value_counts().plot(kind='bar')
```



```
In [52]: sns.barplot(x=df['Cpu brand'],y=df['Price'])
plt.xticks(rotation='vertical')
plt.show()
```



```
In [53]: df.drop(columns=['Cpu', 'Cpu Name'], inplace=True)
```


```
In [54]: df.head()
```

Out[54]:


	Company	TypeName	Ram	Memory	Gpu	OpSys	Weight	Price	Touchscreen	Ips
0	Apple	Ultrabook	8	128GB SSD	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	0	1
1	Apple	Ultrabook	8	128GB Flash Storage	Intel HD Graphics 6000	macOS	1.34	47895.5232	0	0
2	HP	Notebook	8	256GB SSD	Intel HD Graphics 620	No OS	1.86	30636.0000	0	0
3	Apple	Ultrabook	16	512GB SSD	AMD Radeon Pro 455	macOS	1.83	135195.3360	0	1
4	Apple	Ultrabook	8	256GB SSD	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	0	1

```
In [55]: df['Ram'].value_counts().plot(kind='bar')
```

Out[55]: <AxesSubplot:>



```
In [56]: sns.barplot(x=df['Ram'],y=df['Price'])
plt.xticks(rotation='vertical')
plt.show()
```



```
In [57]: df['Memory'].value_counts()
```

Out[57]:

256GB SSD	412
1TB HDD	223
500GB HDD	132
512GB SSD	118
128GB SSD + 1TB HDD	94
128GB SSD	76
256GB SSD + 1TB HDD	73
32GB Flash Storage	38
2TB HDD	16
64GB Flash Storage	15
512GB SSD + 1TB HDD	14
1TB SSD	14
256GB SSD + 2TB HDD	10
1.0TB Hybrid	9
256GB Flash Storage	8
16GB Flash Storage	7
32GB SSD	6
180GB SSD	5
128GB Flash Storage	4
512GB SSD + 2TB HDD	3
16GB SSD	3
512GB Flash Storage	2
128GB SSD + 2TB HDD	2
256GB SSD + 256GB SSD	2

256GB SSD + 500GB HDD	2
1TB SSD + 1TB HDD	2
512GB SSD + 512GB SSD	1
240GB SSD	1
256GB SSD + 1.0TB Hybrid	1
64GB SSD	1
64GB Flash Storage + 1TB HDD	1
1.0TB HDD	1
128GB HDD	1
512GB SSD + 1.0TB Hybrid	1
1TB HDD + 1TB HDD	1
508GB Hybrid	1
8GB SSD	1
512GB SSD + 256GB SSD	1
32GB HDD	1

Name: Memory, dtype: int64

In [58]:

```
df['Memory'] = df['Memory'].astype(str).replace('\.0', '', regex=True)
df["Memory"] = df["Memory"].str.replace('GB', '')
df["Memory"] = df["Memory"].str.replace('TB', '000')
new = df["Memory"].str.split("+", n = 1, expand = True)

df["first"] = new[0]
df["first"] = df["first"].str.strip()

df["second"] = new[1]

df["Layer1HDD"] = df["first"].apply(lambda x: 1 if "HDD" in x else 0)
df["Layer1SSD"] = df["first"].apply(lambda x: 1 if "SSD" in x else 0)
df["Layer1Hybrid"] = df["first"].apply(lambda x: 1 if "Hybrid" in x else 0)
df["Layer1Flash_Storage"] = df["first"].apply(lambda x: 1 if "Flash Storage" in x else 0)

df['first'] = df['first'].str.replace(r'\D', '')

df["second"].fillna("0", inplace = True)

df["Layer2HDD"] = df["second"].apply(lambda x: 1 if "HDD" in x else 0)
df["Layer2SSD"] = df["second"].apply(lambda x: 1 if "SSD" in x else 0)
df["Layer2Hybrid"] = df["second"].apply(lambda x: 1 if "Hybrid" in x else 0)
df["Layer2Flash_Storage"] = df["second"].apply(lambda x: 1 if "Flash Storage" in x else 0)

df['second'] = df['second'].str.replace(r'\D', '')

df["first"] = df["first"].astype(int)
df["second"] = df["second"].astype(int)

df["HDD"] = (df["first"] * df["Layer1HDD"] + df["second"] * df["Layer2HDD"])
df["SSD"] = (df["first"] * df["Layer1SSD"] + df["second"] * df["Layer2SSD"])
df["Hybrid"] = (df["first"] * df["Layer1Hybrid"] + df["second"] * df["Layer2Hybrid"])
df["Flash_Storage"] = (df["first"] * df["Layer1Flash_Storage"] + df["second"] * df["Layer2Flash_Storage"])

df.drop(columns=['first', 'second', 'Layer1HDD', 'Layer1SSD', 'Layer1Hybrid',
                 'Layer1Flash_Storage', 'Layer2HDD', 'Layer2SSD', 'Layer2Hybrid',
                 'Layer2Flash_Storage'], inplace=True)
```

<ipython-input-58-10829db803de>:16: FutureWarning: The default value of regex will change from True to False in a future version.

```
df['first'] = df['first'].str.replace(r'\D', '')
```

<ipython-input-58-10829db803de>:25: FutureWarning: The default value of regex will change from True to False in a future version.

```
df['second'] = df['second'].str.replace(r'\D', '')
```

In [59]:

```
df.sample(5)
```

Out[59]:

	Company	TypeName	Ram	Memory	Gpu	OpSys	Weight	Price	Touchscreen
379	Acer	Notebook	8	1000 HDD	Nvidia GeForce GTX 1050	Linux	2.4	45074.8800	0
907	Acer	Notebook	4	16 Flash Storage	Intel HD Graphics	Chrome OS	2.2	19127.5200	0
541	Dell	Notebook	4	128 SSD	Intel HD Graphics 520	Windows 10	1.6	39207.1536	0
1267	Dell	Notebook	8	1000 HDD	AMD Radeon R5 M430	Linux	2.3	42943.1472	0
1202	Dell	2 in 1 Convertible	16	512 SSD	Intel UHD Graphics 620	Windows 10	2.0	68464.8000	1



In [60]: `df.drop(columns=['Memory'],inplace=True)`

In [61]: `df.head()`

Out[61]:

	Company	TypeName	Ram	Gpu	OpSys	Weight	Price	Touchscreen	lps	pp
0	Apple	Ultrabook	8	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	0	1	226.983005
1	Apple	Ultrabook	8	Intel HD Graphics 6000	macOS	1.34	47895.5232	0	0	127.677940
2	HP	Notebook	8	Intel HD Graphics 620	No OS	1.86	30636.0000	0	0	141.211998
3	Apple	Ultrabook	16	AMD Radeon Pro 455	macOS	1.83	135195.3360	0	1	220.534624
4	Apple	Ultrabook	8	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	0	1	226.983005



In [62]: `df.corr()['Price']`

Out[62]: Ram 0.743007

```

Weight      0.210370
Price       1.000000
Touchscreen 0.191226
Ips         0.252208
ppi         0.473487
HDD         -0.096441
SSD         0.670799
Hybrid      0.007989
Flash_Storage -0.040511
Name: Price, dtype: float64

```

```
In [63]: df.drop(columns=['Hybrid', 'Flash_Storage'], inplace=True)
```

```
In [64]: df.head()
```

```
Out[64]:
```

	Company	TypeName	Ram	Gpu	OpSys	Weight	Price	Touchscreen	Ips	ppi
0	Apple	Ultrabook	8	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	0	1	226.983005
1	Apple	Ultrabook	8	Intel HD Graphics 6000	macOS	1.34	47895.5232	0	0	127.677940
2	HP	Notebook	8	Intel HD Graphics 620	No OS	1.86	30636.0000	0	0	141.211998
3	Apple	Ultrabook	16	AMD Radeon Pro 455	macOS	1.83	135195.3360	0	1	220.534624
4	Apple	Ultrabook	8	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	0	1	226.983005

```
In [65]: df['Gpu'].value_counts()
```

```
Out[65]: Intel HD Graphics 620      281
Intel HD Graphics 520      185
Intel UHD Graphics 620       68
Nvidia GeForce GTX 1050       66
Nvidia GeForce GTX 1060       48
...
AMD Radeon Pro 555           1
Intel Iris Pro Graphics      1
AMD FirePro W5130M           1
AMD Radeon RX 560            1
Nvidia Quadro M500M          1
Name: Gpu, Length: 110, dtype: int64
```

```
In [66]: df['Gpu brand'] = df['Gpu'].apply(lambda x:x.split()[0])
```

```
In [67]: df.head()
```

Out[67]:

	Company	TypeName	Ram	Gpu	OpSys	Weight	Price	Touchscreen	Ips	pp
0	Apple	Ultrabook	8	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832	0	1	226.983005
1	Apple	Ultrabook	8	Intel HD Graphics 6000	macOS	1.34	47895.5232	0	0	127.677940
2	HP	Notebook	8	Intel HD Graphics 620	No OS	1.86	30636.0000	0	0	141.211998
3	Apple	Ultrabook	16	AMD Radeon Pro 455	macOS	1.83	135195.3360	0	1	220.534624
4	Apple	Ultrabook	8	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080	0	1	226.983005



In [68]:

```
df['Gpu brand'].value_counts()
```

Out[68]:

```
Intel      722
Nvidia    400
AMD        180
ARM         1
Name: Gpu brand, dtype: int64
```

In [69]:

```
df = df[df['Gpu brand'] != 'ARM']
```

In [70]:


```
df['Gpu brand'].value_counts()
```

Out[70]:

```
Intel      722
Nvidia    400
AMD        180
Name: Gpu brand, dtype: int64
```

In [71]:

```
sns.barplot(x=df['Gpu brand'],y=df['Price'],estimator=np.median)
plt.xticks(rotation='vertical')
plt.show()
```



```
In [72]: df.drop(columns=['Gpu'],inplace=True)
```

```
In [73]: df.head()
```

```
Out[73]:
```


	Company	TypeName	Ram	OpSys	Weight	Price	Touchscreen	Ips	ppi	Cpu brand
0	Apple	Ultrabook	8	macOS	1.37	71378.6832	0	1	226.983005	Intel Core i5
1	Apple	Ultrabook	8	macOS	1.34	47895.5232	0	0	127.677940	Intel Core i5
2	HP	Notebook	8	No OS	1.86	30636.0000	0	0	141.211998	Intel Core i5
3	Apple	Ultrabook	16	macOS	1.83	135195.3360	0	1	220.534624	Intel Core i7
4	Apple	Ultrabook	8	macOS	1.37	96095.8080	0	1	226.983005	Intel Core i5

```
In [74]: df['OpSys'].value_counts()
```

```
Out[74]: Windows 10      1072
No OS                66
Linux                62
Windows 7            45
Chrome OS            26
macOS                13
Mac OS X             8
Windows 10 S         8
Android              2
Name: OpSys, dtype: int64
```

```
In [75]: sns.barplot(x=df['OpSys'],y=df['Price'])
```

```
plt.xticks(rotation='vertical')
plt.show()
```



```
In [76]: def cat_os(inp):
          if inp == 'Windows 10' or inp == 'Windows 7' or inp == 'Windows 10 S':
              return 'Windows'
          elif inp == 'macOS' or inp == 'Mac OS X':
              return 'Mac'
          else:
              return 'Others/No OS/Linux'
```

```
In [77]: df['os'] = df['OpSys'].apply(cat_os)
```


```
In [78]: df.head()
```

```
Out[78]:
```

	Company	TypeName	Ram	OpSys	Weight	Price	Touchscreen	Ips	ppi	Cpu brand
0	Apple	Ultrabook	8	macOS	1.37	71378.6832	0	1	226.983005	Intel Core i5
1	Apple	Ultrabook	8	macOS	1.34	47895.5232	0	0	127.677940	Intel Core i5
2	HP	Notebook	8	No OS	1.86	30636.0000	0	0	141.211998	Intel Core i5
3	Apple	Ultrabook	16	macOS	1.83	135195.3360	0	1	220.534624	Intel Core i7
4	Apple	Ultrabook	8	macOS	1.37	96095.8080	0	1	226.983005	Intel Core i5

```
In [79]: df.drop(columns=['OpSys'],inplace=True)
```


```
In [80]: sns.barplot(x=df['os'],y=df['Price'])  
plt.xticks(rotation='vertical')  
plt.show()
```



```
In [81]: sns.distplot(df['Weight'])
```


C:\Users\Varun kumar\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[81]: <AxesSubplot:xlabel='Weight', ylabel='Density'>
```



```
In [82]: sns.scatterplot(x=df['Weight'],y=df['Price'])
```

```
Out[82]: <AxesSubplot:xlabel='Weight', ylabel='Price'>
```

In [83]: `df.corr()['Price']`


Out[83]:

Ram	0.742905
Weight	0.209867
Price	1.000000
Touchscreen	0.192917
Ips	0.253320
ppi	0.475368
HDD	-0.096891
SSD	0.670660

Name: Price, dtype: float64

In [140... `sns.heatmap(df.corr(),annot=True)`


Out[140... <AxesSubplot:>



In [85]: `sns.distplot(np.log(df['Price']))`

C:\Users\Varun kumar\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

Out[85]: <AxesSubplot:xlabel='Price', ylabel='Density'>



```
In [86]: X = df.drop(columns=['Price'])
        y = np.log(df['Price'])
```

```
In [87]: X
```

Out[87]:

	Company	TypeName	Ram	Weight	Touchscreen	Ips	ppi	Cpu brand	HDD	SSD	C	br
0	Apple	Ultrabook	8	1.37	0	1	226.983005	Intel Core i5	0	128	I	
1	Apple	Ultrabook	8	1.34	0	0	127.677940	Intel Core i5	0	0	I	
2	HP	Notebook	8	1.86	0	0	141.211998	Intel Core i5	0	256	I	
3	Apple	Ultrabook	16	1.83	0	1	220.534624	Intel Core i7	0	512	A	
4	Apple	Ultrabook	8	1.37	0	1	226.983005	Intel Core i5	0	256	I	
...		
1298	Lenovo	2 in 1 Convertible	4	1.80	1	1	157.350512	Intel Core i7	0	128	I	
1299	Lenovo	2 in 1 Convertible	16	1.30	1	1	276.053530	Intel Core i7	0	512	I	
1300	Lenovo	Notebook	2	1.50	0	0	111.935204	Other Intel Processor	0	0	I	
1301	HP	Notebook	6	2.19	0	0	100.454670	Intel Core i7	1000	0	A	
1302	Asus	Notebook	4	2.20	0	0	100.454670	Other Intel Processor	500	0	I	

1302 rows × 12 columns



In [88]:

y

Out[88]:

0

11.175755

1

10.776777

2

10.329931

3

11.814476

4

11.473101

...

1298

10.433899

1299

11.288115

1300

9.409283

1301

10.614129

1302

9.886358

Name: Price, Length: 1302, dtype: float64

In [89]:

from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.15,random_state=2)

In [90]:

X_train

Out[90]:

	Company	TypeName	Ram	Weight	Touchscreen	Ips	ppi	Cpu brand	HDD	SSD	C	br
183	Toshiba	Notebook	8	2.00	0	0	100.454670	Intel Core i5	0	128	I	
1141	MSI	Gaming	8	2.40	0	0	141.211998	Intel Core i7	1000	128	Nv	
1049	Asus	Netbook	4	1.20	0	0	135.094211	Other Intel Processor	0	0	I	
1020	Dell	2 in 1 Convertible	4	2.08	1	1	141.211998	Intel Core i3	1000	0	I	
878	Dell	Notebook	4	2.18	0	0	141.211998	Intel Core i5	1000	128	Nv	
...		
466	Acer	Notebook	4	2.20	0	0	100.454670	Intel Core i3	500	0	Nv	
299	Asus	Ultrabook	16	1.63	0	0	141.211998	Intel Core i7	0	512	Nv	
493	Acer	Notebook	8	2.20	0	0	100.454670	AMD Processor	1000	0	A	
527	Lenovo	Notebook	8	2.20	0	0	100.454670	Intel Core i3	2000	0	Nv	
1193	Apple	Ultrabook	8	0.92	0	1	226.415547	Other Intel Processor	0	0	I	

1106 rows × 12 columns

In [91]:

from sklearn.compose import ColumnTransformer

from sklearn.pipeline import Pipeline

```
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import r2_score, mean_absolute_error
```

In [92]:

```
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, AdaBoostClassifier
from sklearn.svm import SVR
from xgboost import XGBRegressor
```

Linear regression

In [93]:

```
step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0, 1, 7, 10, 11])
], remainder='passthrough')

step2 = LinearRegression()

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

R2 score 0.8073277448418522
MAE 0.21017827976429174

Ridge Regression

In [94]:

```
step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0, 1, 7, 10, 11])
], remainder='passthrough')

step2 = Ridge(alpha=10)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

R2 score 0.8127331031311809
MAE 0.20926802242582976

Lasso Regression

In [139]...

```
step1 = ColumnTransformer(transformers=[
```

```
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0, 1, 7, 10, 11])
], remainder='passthrough')

step2 = Lasso(alpha=0.002)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

R2 score 0.8048986293268664
MAE 0.21252926520036242

KNN

Decision Tree

In [97]:

```
step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0, 1, 7, 10, 11])
], remainder='passthrough')

step2 = DecisionTreeRegressor(max_depth=8)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

R2 score 0.831356290461679
MAE 0.18686616252396962

SVM

In [122...]

```
step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0, 1, 7, 10, 11])
], remainder='passthrough')

step2 = SVR(kernel='rbf', C=10000, epsilon=0.1)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)
```

```
print('R2 score',r2_score(y_test,y_pred))
print('MAE',mean_absolute_error(y_test,y_pred))
```

R2 score 0.8083180902257614
MAE 0.20239059427481307

Random Forest

In [130...

```
step1 = ColumnTransformer(transformers=[
    ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])
],remainder='passthrough')

step2 = RandomForestRegressor(n_estimators=800,
                              random_state=3,
                              max_samples=0.5,
                              max_features=0.85,
                              max_depth=19)

pipe = Pipeline([
    ('step1',step1),
    ('step2',step2)
])

pipe.fit(X_train,y_train)

y_pred = pipe.predict(X_test)

print('R2 score',r2_score(y_test,y_pred))
print('MAE',mean_absolute_error(y_test,y_pred))
```

R2 score 0.8842238734380571
MAE 0.16010485229215574

ExtraTrees

AdaBoost

In [114...

```
step1 = ColumnTransformer(transformers=[
    ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])
],remainder='passthrough')

step2 = AdaBoostRegressor(n_estimators=80,learning_rate=1.1)

pipe = Pipeline([
    ('step1',step1),
    ('step2',step2)
])

pipe.fit(X_train,y_train)

y_pred = pipe.predict(X_test)

print('R2 score',r2_score(y_test,y_pred))
print('MAE',mean_absolute_error(y_test,y_pred))
```

R2 score 0.7933259540611934
MAE 0.22899914522846598

Gradient Boost

```
In [131... step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0,1,7,10,11])
], remainder='passthrough')

step2 = GradientBoostingRegressor(n_estimators=500)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

R2 score 0.8803238273292837
MAE 0.16019652415915453

XgBoost

```
In [137... step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0,1,7,10,11])
], remainder='passthrough')

step2 = XGBRegressor(n_estimators=140, max_depth=5, learning_rate=0.5)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))
```

R2 score 0.8746950689623854
MAE 0.16582715402265588

Voting Regressor

```
In [108... from sklearn.ensemble import VotingRegressor, StackingRegressor

step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0,1,7,10,11])
], remainder='passthrough')

estimators = [
    ('rf', RandomForestRegressor(n_estimators=350, random_state=3, max_samples=0.5, max_depth=5)),
    ('gbdt', GradientBoostingRegressor(n_estimators=100, max_features=0.5)),
    ('xgb', XGBRegressor(n_estimators=25, learning_rate=0.3, max_depth=5))
]

step2 = StackingRegressor(estimators=estimators, final_estimator=Ridge(alpha=100))

pipe = Pipeline([
```

```
    ('step1', step1),  
    ('step2', step2)  
])  
  
pipe.fit(X_train, y_train)  
  
y_pred = pipe.predict(X_test)  
  
print('R2 score', r2_score(y_test, y_pred))  
print('MAE', mean_absolute_error(y_test, y_pred))
```

R2 score 0.8809533645736113
MAE 0.16662383385951088

In []:

In []:

In []: